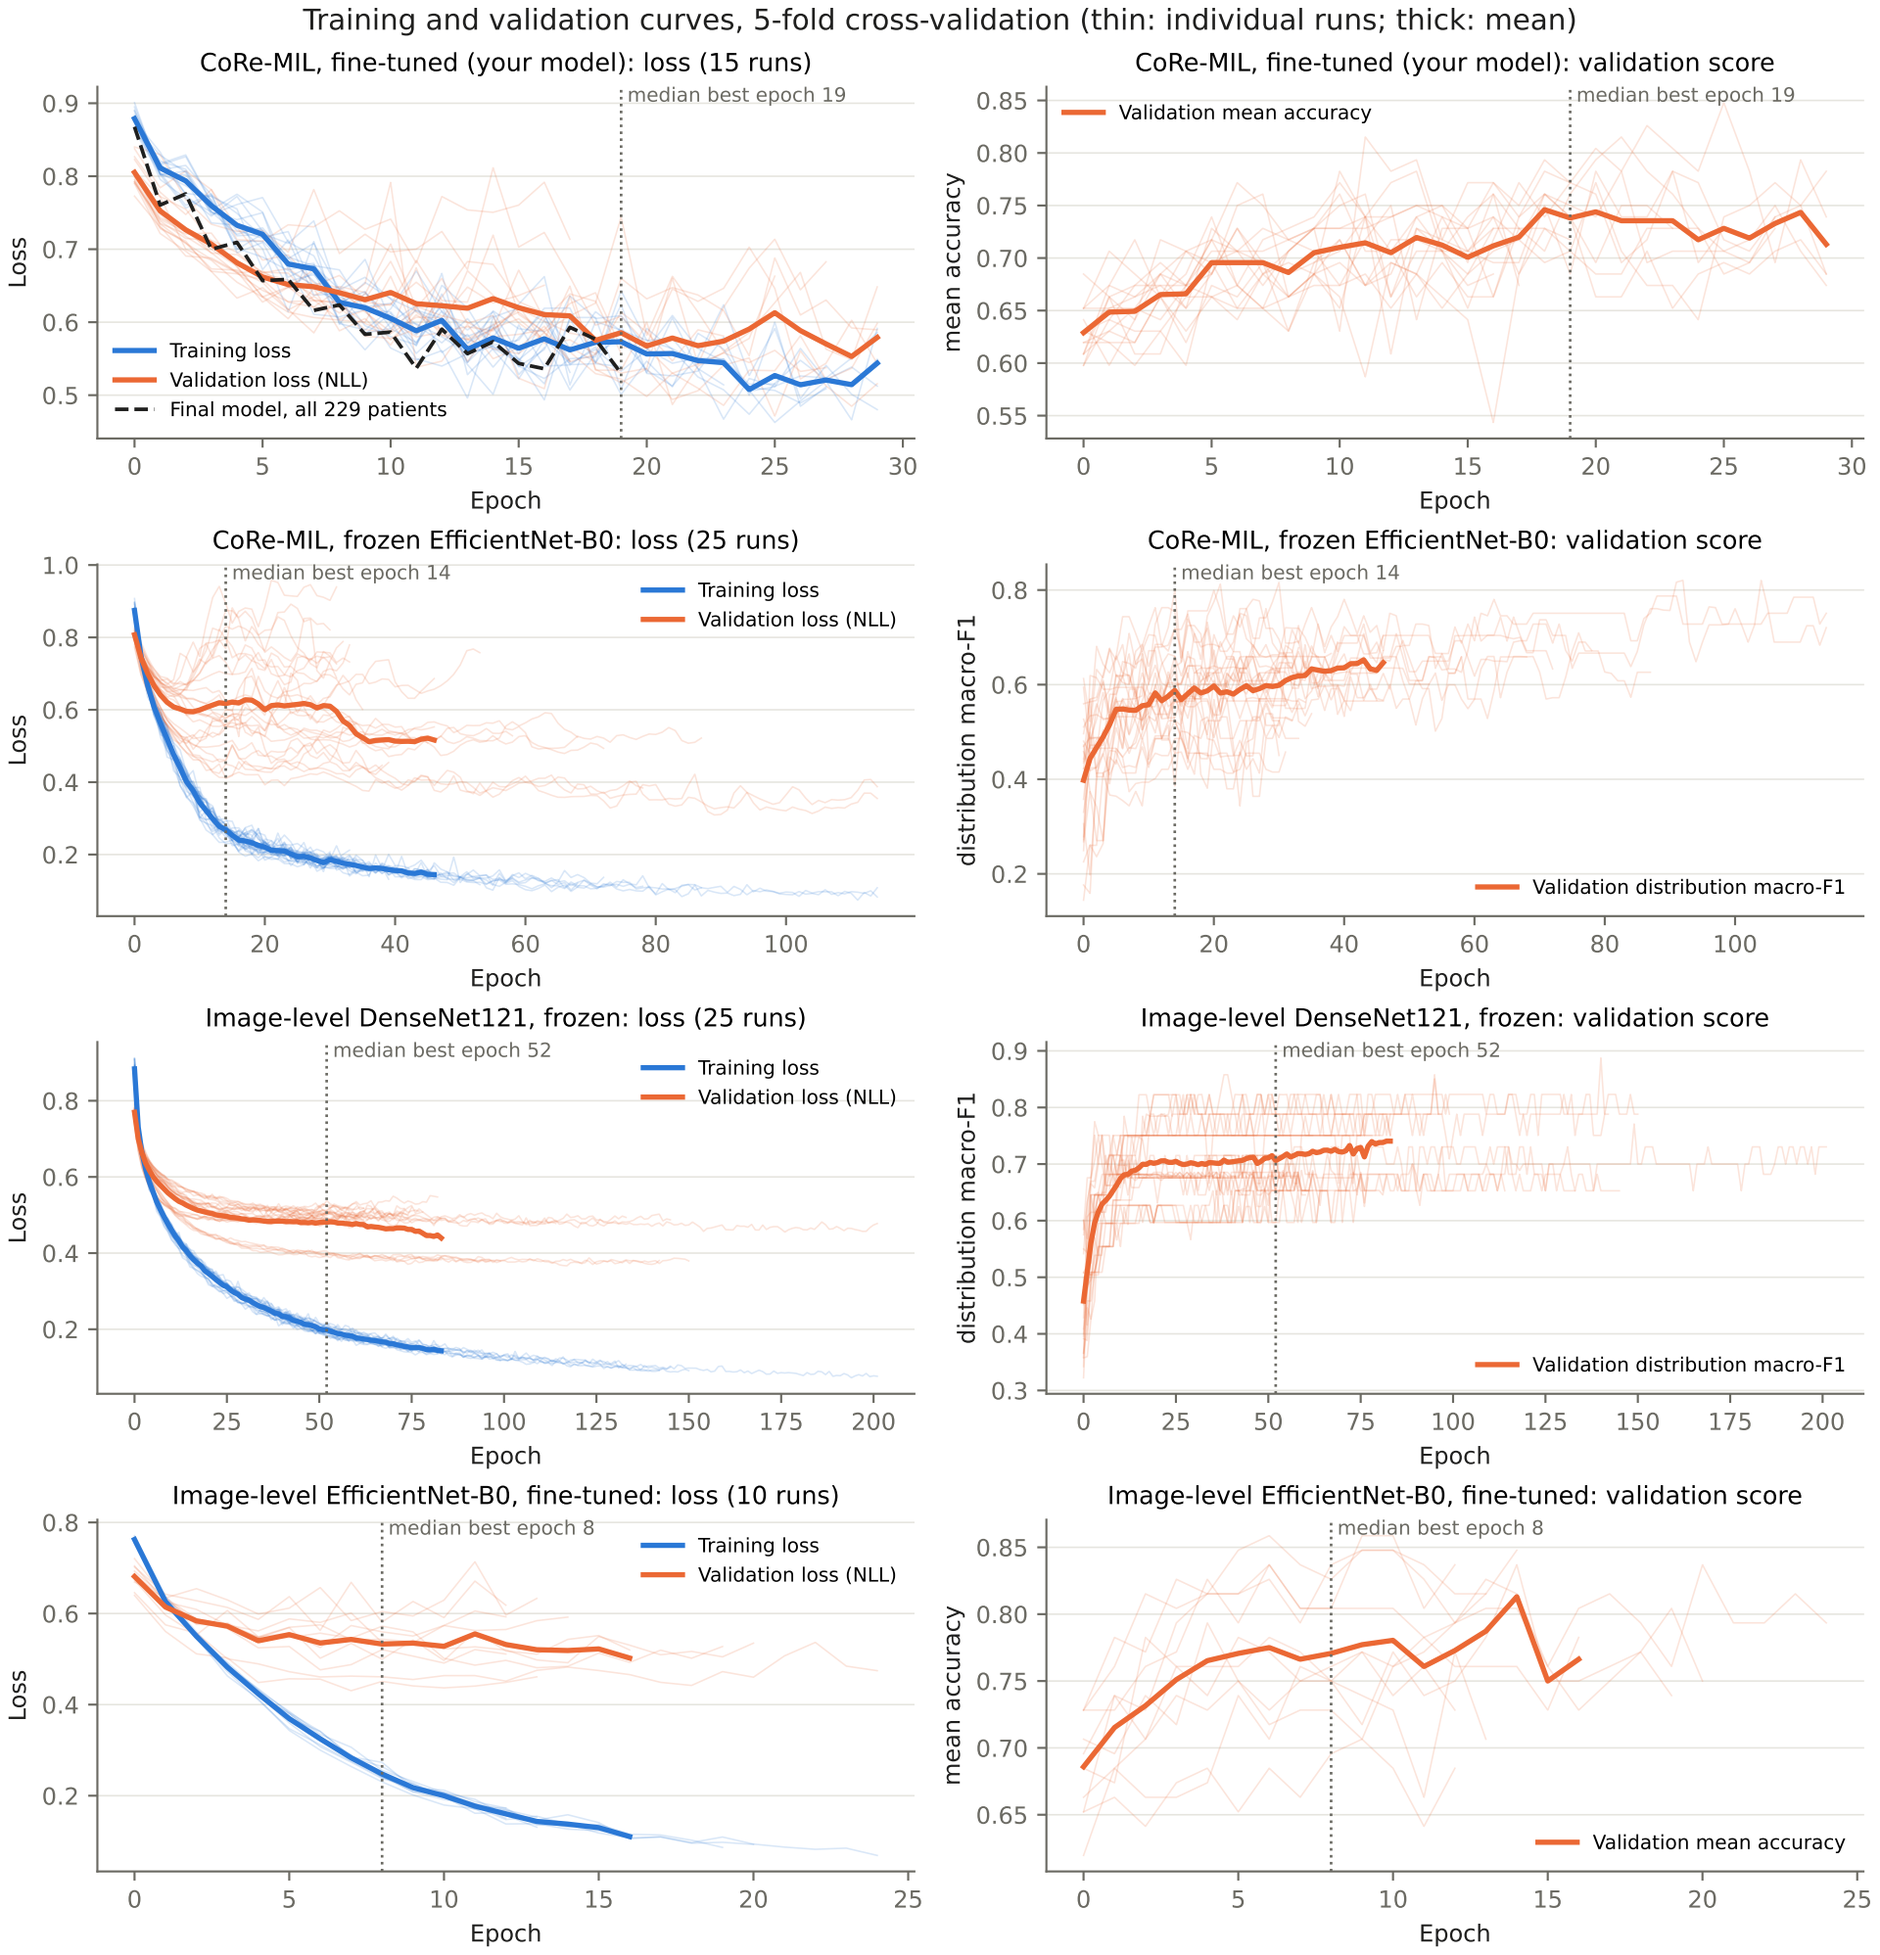

Training and validation curves, 5-fold cross-validation (thin: individual runs; thick: mean)
CoRe-MIL, fine-tuned (your model): loss (15 runs)
median best epoch 19
Training loss
Validation loss (NLL)
Final model, all 229 patients
Loss
Epoch
CoRe-MIL, fine-tuned (your model): validation score
Validation mean accuracy
median best epoch 19
mean accuracy
Epoch
CoRe-MIL, frozen EfficientNet-B0: loss (25 runs)
median best epoch 14
Training loss
Validation loss (NLL)
Loss
Epoch
CoRe-MIL, frozen EfficientNet-B0: validation score
median best epoch 14
Validation distribution macro-F1
distribution macro-F1
Epoch
Image-level DenseNet121, frozen: loss (25 runs)
median best epoch 52
Training loss
Validation loss (NLL)
Loss
Epoch
Image-level DenseNet121, frozen: validation score
median best epoch 52
Validation distribution macro-F1
distribution macro-F1
Epoch
Image-level EfficientNet-B0, fine-tuned: loss (10 runs)
median best epoch 8
Training loss
Validation loss (NLL)
Loss
Epoch
Image-level EfficientNet-B0, fine-tuned: validation score
median best epoch 8
Validation mean accuracy
mean accuracy
Epoch